



Making Educated Guesses on Password Storage Using Timing Tests

Monika Morrow

Cindee Tran

Introductions

- Monika Morrow
 - Senior Security Consultant at AppSec Consulting
 - <https://github.com/monikamorrow>
 - @FortyTwoWho
- Cindee Tran
 - Senior Security Consultant at AppSec Consulting



Why?

- Detect insecure password storage without code review
- Allow external validation of your secure password storage

Abstract:

It is becoming more popular to limit password lengths to prevent denial of service by submissions of long passwords to computationally expensive password hashing algorithms like PBKDF2, bcrypt, and scrypt. How expensive is it really? Can we use denial of service tactics to learn about an applications password storage mechanism without a code review? Let's find out!

The birth of an idea:

The original problem: What is a reasonable maximum length password to recommend? This unearthed an interesting fact, Django limited its password length to 4096 to prevent Denial of Service attacks against its login which used PBKDF2.

Light bulb! If a system using a secure password storage mechanism can be DoSed because of the computationally expensive algorithms it should be possible to determine if a system is using secure password storage functions by observing the amount of time the server takes to respond to password requests.

<http://arstechnica.com/security/2013/09/long-passwords-are-good-but-too-much-length-can-be-bad-for-security/>

Passwords Then... (we wish!)

- Simple hash
 - MD5
 - SHA256
 - Downfall: Lookup tables
- Salted hash
 - Same algorithms as before
 - Solved: Lookup tables using salt
 - Downfall: Computing power overwhelming

Passwords Now (we hope!)

- Work factor hashing algorithms
 - bcrypt
 - scrypt
 - PBKDF2
 - Solved: Computing power overwhelming using configurable work factor

bcrypt

maximum input length of 56 or 72 bytes

blowfish, expensive keysetup phase

scrypt

Relies on PBKDF2

PBKDF2

computationally expensive key derivation function

Why a Work Factor?

- Hashing is FAST, measured in Mh/s
- SHA256 is slower than MD5 but PC3 still performs 14,416,000 hashes per second

Performance

PC1: Windows 7, 32 bit · Catalyst 14.9 · 1x AMD hd7970 · 1000mhz core clock · oclHashcat v1.35

PC2: Windows 7, 64 bit · ForceWare 347.52 · 1x NVidia gtx580 · stock core clock · oclHashcat v1.35

PC3: Ubuntu 14.04, 64 bit · ForceWare 346.29 · 8x NVidia Titan X · stock core clock · oclHashcat v1.36

PC4: Ubuntu 14.04, 64 bit · Catalyst 14.9 · 8x AMD R9 290X · stock core clock · oclHashcat v1.35

Hash Type	PC1	PC2	PC3	PC4
MD5	8581 Mh/s	2753 Mh/s	115840 Mh/s	92672 Mh/s
SHA1	3037 Mh/s	655 Mh/s	37336 Mh/s	31552 Mh/s
SHA256	1122 Mh/s	355 Mh/s	14416 Mh/s	12288 Mh/s
SHA512	414 Mh/s	104 Mh/s	4976 Mh/s	4552 Mh/s

Credit: <https://hashcat.net/oclhashcat/>

Plaintext link:

<https://hashcat.net/oclhashcat/>

Data appears to be circa early 2015

Our Hypothesis

- The response delay introduced by the password hashing algorithm's work factor can be detected remotely.

What we expected to learn:

What approximate length password would reveal the likely underlying password storage mechanism?

Server Setup (WordPress)

- MD5 (Default)
- bcrypt (wp-bcrypt Plugin)
 - Work Factor 10
 - Work Factor 15
 - Work Factor 20
- scrypt (Encrypt Plugin)
- PBKDF2 (ballast-security-securing-hashing Plugin)
 - Ballast Security's modified PBKDF2 with 100000 iterations
 - Ballast Security original ARC4PBKDF2 with 100000 iterations
 - Classic PBKDF2 with 100000 iterations

To test our hypothesis, we set up an dedicated server hosting environment with 4 WordPress installations – one for each password hashing algorithm. For those that are unfamiliar with WordPress, it's just a simple open source CMS that allow you to easily install plugins so that you can extend or add functionality to the website.

By default for password storage, WordPress adds salt to the password and hashes it with 8 passes of MD5. For bcrypt, scrypt, and PBKDF2, we installed the appropriate plugins from their Plugin Directory and modified the code or settings to use different work factors. For bcrypt, we used three different work factors and for PBKDF2, we tested three different implementations of it.

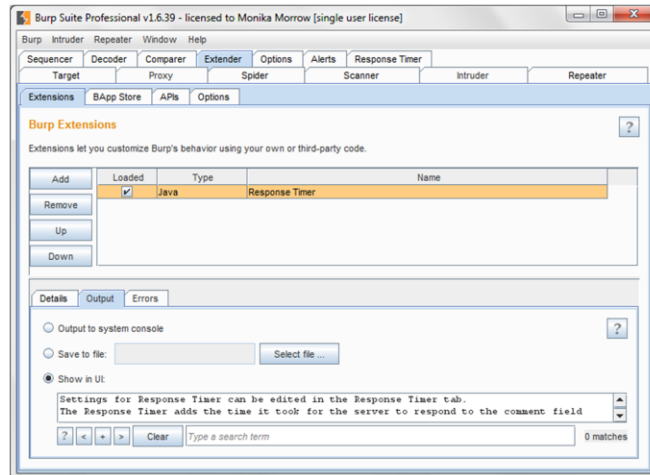
Burp Extension

- Response Timer Extension
 - Average
 - Maximum
 - Minimum
- Data points
 - 100 requests per data point
 - 8 16 24 32 40 48 64 112 160
 - 208 256 4096 65536 1048576

In order to collect all the data for our research, Monika created an awesome Burp extension that captures and auto-populates the average, maximum, and minimum response times.

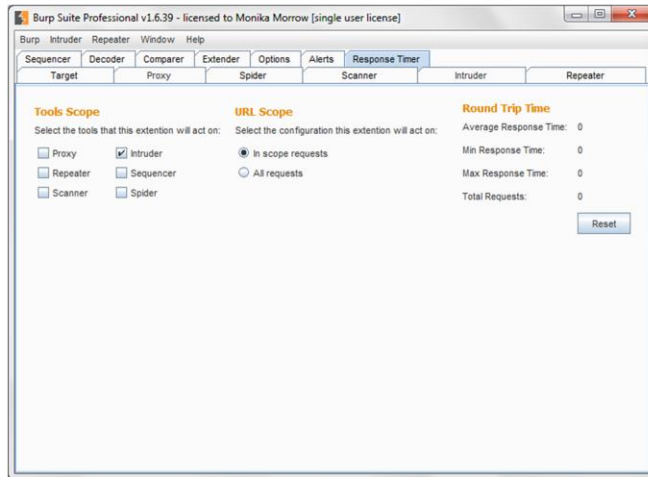
The data points represent the number of random characters we used for each password. For every data point, we tested for each of the 4 different password hashing algorithms and the different configurations for bcrypt and PBKDF2.

The Setup – Load Extension



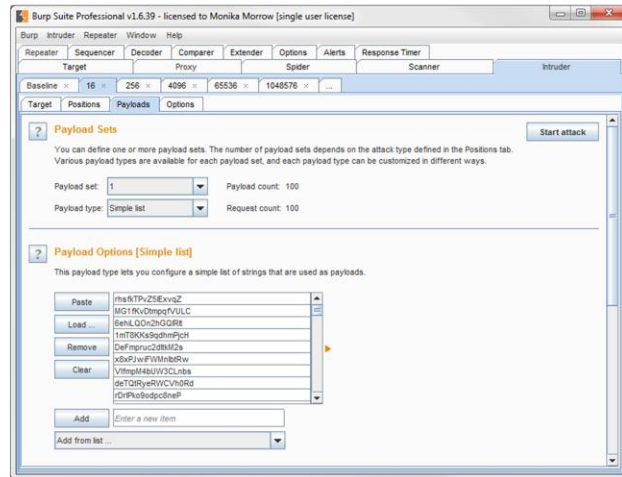
The Burp extension setup is simple, you just load the extension in Burp Extender.

Configure Extension



To configure the burp extension, you need to keep Intruder checked under Tools Scope and also add the domain to scope under the Target tab or just select all requests.

Configure Intruder - Payloads

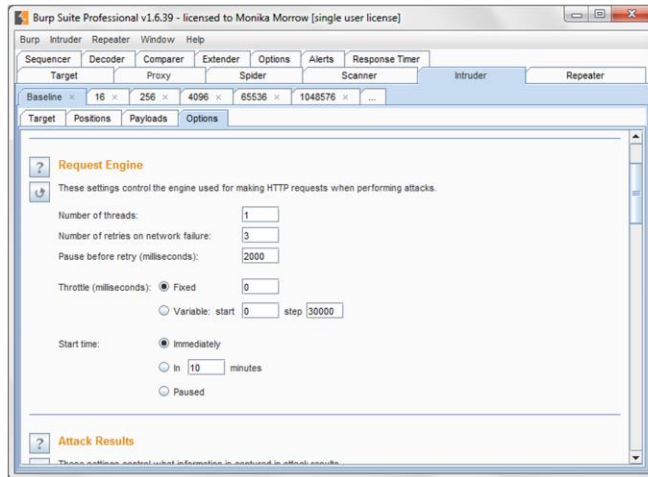


To configure Intruder, we generated random characters for each data point and configured Intruder to use a Simple List payload. The example here is for 16 characters.

Example code for making random passwords

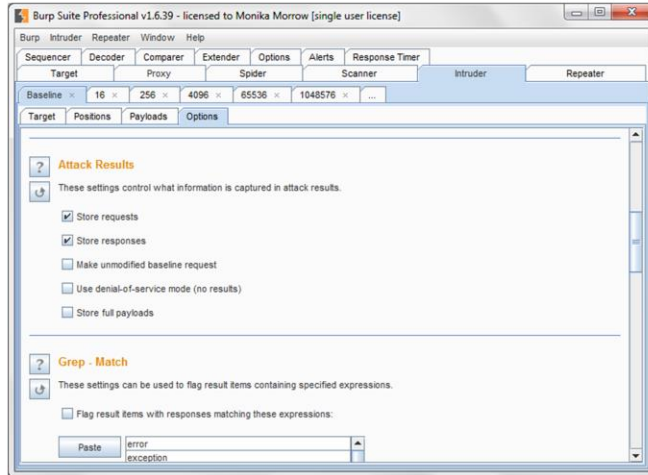
```
#!/bin/bash
COUNTER=0
while [ $COUNTER -lt 20 ]; do
    cat /dev/urandom | tr -dc A-Za-z0-9 | head -c 1048576 >> 1m.txt
    echo \n >> 1m.txt
    let COUNTER=COUNTER+1
done
```

Configure Intruder – Single Thread



In order to maximize our chances of getting valid results, we configured Intruder to run on single thread mode.

Configure Intruder – No Base Request



Another configuration we made was to ensure that the unmodified baseline request was unchecked to keep it from making a request with a shorter password and therefore adding invalid data to the data set.

Data?

“I do know something,
just not with any certainty.”

“I'm a scientist. Certainty is a big word.
I need time to collect my data.”

-Dr. Amy Barnes
Volcano (1997)

Data



Plaintext link:

<https://github.com/MonikaMorrow/PasswordTimingTest/DataSet2016-03-26.xlsx>

What We Learned

- We have more to learn
- Large values don't have as much of an effect as we expected
- Results suggest measurable difference at small values

What We Learned

- bcrypt only uses first 56 or 72 chars
 - Some applications hash passwords before performing bcrypt to ensure all characters are accounted for
- Ensuring limited maximum length password processed by PBKDF2 appears to limit DoS Risk

Call For Moar Data!

- Reproduce our results
- Submit data
 - <https://github.com/MonikaMorrow/PasswordTimingTest>
- Contact us
 - MonikaMorrow@gmail.com
 - Tran.Cindee@gmail.com

Plaintext link:

<https://github.com/MonikaMorrow/PasswordTimingTest>

Q&A

